

1 Three Address Code

1.1 IR for Assembly Code Generation

This is low level IR and is meant for assembly code generation. Typically this IR is present in compiler tool chains of almost all languages. IR resembles high-level assembly. Some properties for such IR are:

- They have register names similar to assembly language. Unlike assembly code, they have unlimited numbers of registers.
- They have control structures like conditional and unconditional branches(similar to assembly). They are lower level than if-else(in C) etc.
- They have opcodes. They have higher level opcodes, something like doing vector addition. These opcodes are very similar to the ones in assembly which make their conversion to assembly opcodes easier. Sometimes IR opcodes could be lower level as well. For eg., some assembly opcodes can do multiple operations (*add*, *multiply*) in one instruction which might not be the case with the IR opcode. These need to be kept into mind while converting from IR to assembly code. An opcode is higher level if its translation to assembly code results in more than 1 instruction. If more than one opcodes translate to one instruction in assembly code, then it is considered lower level opcode.

1.2 IR Characteristics

Each instruction is of the form

$$x = y \text{ op } z \text{ (binary operation)}$$

or

$$x = \text{op } y \text{ (unary operation)}$$

where y and z are registers or constants, can be scalars or vectors. This type of IR is called *Three-Address Code*. This is because, for each operation, there are at max 3 addresses (x, y, z). Addresses are registers or memory addresses.

1.3 Three Address Code Property

Expression $x + y * z$ translates to

$$t1 = y * z$$

$$t2 = x + t1$$

Property : Each sub-expression has a name which makes compilation easier later on from the perspective of optimization. Assembly operations such as *load* and *store* are also supported.

1.4 IR Code Generation

This is very similar to assembly code generation.

Define function $igen(e, t)$ as a function which

(INPUT) : takes an expression e and a value t . t represents the register name in which the value of expression should be computed.

(OUTPUT) : $igen(e, t)$ generates code to compute the value of expression e in register t . (Currently assuming that there are unlimited number of registers)

1.4.1 IR CodeGen Example

$$\begin{aligned} igen(e_1 + e_2, t) = & \\ & igen(e_1, t_1) \\ & igen(e_2, t_2) \\ & t = t_1 + t_2 \end{aligned}$$

t_1 and t_2 are fresh registers.

1.4.2 Three Address Code IR

LLVM (*Low level Virtual Machine*) is the one of most popular *Three Address Code* IRs.

Types in LLVM are :

(integer) $i1, i8, i32..$ ($i1$ stands for 1-bit integer and so on..)

(float)

(pointers) $i32^*, i32^{**}..$

(vector types) $\langle 4 * i32 \rangle$

(struct types) $\{i32, i16\}$

The size of the integer is not consequential to optimization, so we make this decision by the time we reach IR.

Risc-like opcodes : *LLVM* has opcodes like add, load, store, call, ret, branch, conditional branch... Unlike assembly code, these opcodes can have unbounded number of arguments.

2 Phi Nodes

In SSA IR, we encountered the issue of figuring out which variable version to use after a join point where multiple paths were coming in with different versions of the same variable. As a solution to this, phi nodes or phi functions are extremely helpful that could be placed only at the beginning of a basic block with multiple edges coming in (join point).

These phi nodes can be placed at each join point for every variable in the program. This placement strategy is a bit wasteful as illustrated in the following diagram:



Figure 1: Placement of Φ nodes

2.1 Path Convergence Criterion

Consider z to be node with multiple edges coming in (thus it is a join point). We need Φ node for variable a at node z if and only if

1. There is a block x containing the definition of a .
2. There is a block y (not x) containing the definition of a .
3. There are non empty paths P_{xz} and P_{yz} from x to z and y to z respectively.




Figure 2: CFG for condition 1, 2, 3

4. P_{xz} and P_{yz} should not have any node in common except z .
5. The node z does not appear within both P_{xz} and P_{yz} , prior to the end. It is fine if z appears only in one of the path before end.

2.2 Iterative Fixed Point Algorithm

while {there are nodes x, y, z satisfying condition 1-5 and z does not contain a Φ node for a }
do {insert $a \leftarrow \Phi(a_1, a_2, \dots, a_j)$ }

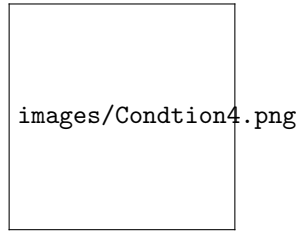


Figure 3: Violation of Condition 4




Figure 4: Condition 5: Left CFG does not need Φ node, right CFG needs Φ node

The Φ function has as many 'a' arguments as there are predecessors of z. Since the conditions 1-5 are both sufficient and necessary, the above algorithm is sound and complete.

3 Optimization Overview

Most intermediate representations are organised as *control flow graphs (CFG)* over *basic blocks*.

3.1 Basic Blocks

A basic block is a single entry, single exit, straight line code segment. More formally, it is a maximal sequence of instructions with no labels (except at the first instruction) and no jumps (except at the last instruction).

In the following figure, (a) is a basic block. (b) is not a basic block because it has multiple exits; (c) is not a basic block because it has multiple entries; (d) is not a basic block because it does not represent a straight line code.

Consider the following example of a single basic block:

1. L:
 2. $t := 2 * x$
 3. $w := t + x$
 4. if $w > 3$ goto L'
- Is it ok to change (3) to $w := 3 * x$? It is ok. If addition of two numbers is more expensive as compared to multiplication of a number with a constant, the above change can be considered an optimization.




Figure 5: Examples and Counterexamples for basic block

- Is it ok to change (4) to `if x > 1 goto L'` ? It is not ok. Because here `x` is a finite bounded integer and due to overflow, it is possible to have $3 * x > 3$ and not `x > 1`.
- Is it ok to remove (2) ? Depends on whether variable `t` is used later in the code or not.

3.2 Control Flow Graph

A control flow graph is a directed graph with

- basic blocks as nodes
- edge from block A to block B if the execution can pass from the last instruction in A to the first instruction in B, for example
 - the last instruction in A is: `jump Lb`
 - the last instruction in A is: `if id1 = id2 then goto Lb`
 - execution can fall-through from block A to block B

Consider the following example of a CFG:

The body of a method (or procedure) can be represented as a control-flow graph. There is one initial node (entry node). All "return" nodes are terminal.

3.3 Optimization

The optimizations are performed on the control flow graph of the intermediate representation of the code to improve program's resource utilization:

- Execution Time
- Code Size



Figure 6: CFG Example

- Memory Usage
- Frequency of Disk I/O operations (or network operations)
- Power Consumption (Not same as energy)

It is important to remember that optimization should not alter the meaning of the program. It should not alter what the program computes.

For example, the following optimization changes the meaning of the program and thus, it is not a valid optimization.

$w := 3 * x$ if $w > 3$ goto L
cannot be converted to: if $x > 1$ goto L

3.4 Typical Granularity of Optimization

1. Local Optimizations

- applied to a basic block in isolation
- easiest to implement

2. Global Optimizations

- applied to a CFG (method body) in isolation while crossing the boundaries of basic blocks

3. Inter-Procedural Optimizations

- applied across method (CFG) boundaries
- difficult to implement but usually most effective

3.5 Economics of the Optimization

Optimizations are more of an art rather than science. The current state of the art methods are based on the concept of "Maximum benefit for minimum cost" where cost can denote the development and integration costs of the optimization.

- Some optimizations are hard to implement
- Some optimizations require large compilation time
- Some optimizations have low payoff (the benefits) and it is often difficult to quantify payoff.

4 Local Optimizations

Local Optimizations are the simplest form of optimizations that can be performed by considering a single basic block in isolation. Some possible examples are:

- **Elimination of No-ops:** Some statements can be deleted. $x := x + 0$, $x := x * 1$, $x := x|0$
- **Algebraic Simplification:** Some statements can be simplified

$$x := x * 0 \rightarrow x := 0$$

$$y := x * *2 \rightarrow y := x * x$$

$$x := x * 8 \rightarrow x := x << 3$$

$$x := x * 15 \rightarrow t := x << 4; x := t - x$$

The above code replacements are meaningful optimizations only if the RHS code would perform better than LHS code which also depends on the underlying hardware.

- **Constant Propagation or Constant Folding:** For statement $x := y \text{ op } z$ where y and z are constants, this statement can be computed at compile time.

$$x := 2 + 2 \rightarrow x := 4$$

$$\text{if } 2 < 0 \text{ jump } L \rightarrow \text{No-op}$$

$$\text{if } 2 > 0 \text{ jump } L \rightarrow \text{jump } L$$

More specifically, constant folding is about performing computations at compile time and constant propagation is about propagating constants throughout the program.

4.1 Dead Code Elimination

Dead code is the code that is unreachable from the initial block, for example, CFG node with no incoming edges. Removing unreachable code makes the program smaller (and sometimes faster due to fewer cache misses).

Why would unreachable block occur?

```
#define DEBUG 0
...
if (DEBUG) {
...
}
```

In the above example, after performing constant propagation, `DEBUG` would be replaced with `0`. Further constant folding would remove the jump statement to the 'if' block and it would be converted to dead code which can be eliminated.

Unreachable block may also occur when libraries are imported. Libraries might have a large number of functions but usually only a small fraction of functions are actually used. Then the code for other functions constitute dead code. Usually other optimizations may result in more dead code.

4.2 Common Subexpression Elimination

Consider the case of following code with the assumption that the value of `x`, `y`, and `z` remain unchanged after the first statement, then we can perform the following replacement to use the pre-computed value.

$$x := y + z; \dots; \dots; w := y + z; \rightarrow x := y + z; \dots; \dots; w := x;$$

In this case, SSA IR particularly helps in doing away with the initial assumption since SSA IR enforces the property that each variable can be assigned only once. In short, SSA IR makes more values available simultaneously.

$$x_1 := y + z; x_2 := x_1 + 3; w := y + z \rightarrow x_1 := y + z; x_2 := x_1 + 3; w := x_1$$

In the above example, the common subexpression being eliminated is `y + z`. The three address code also makes it easier to identify the uses of a subexpression in the code.

4.3 Copy Propagation

If we see `w := x`, replace subsequent uses of `w` with `x` (and eliminate this statement)

Before: `x := y + z; w := x + 3; v := x; u := v + 3;`

After: `x := y + z; w := x + 3; u := x + 3;`

This optimization leads to less number of copies leading to need of fewer registers. It also activates further optimizations like common subexpression elimination.

4.4 Example 1

Original Code: $x := y + z$; $w := x + 3$; $v := y + z$; $u := v + 3$;

After CSE: $x := y + z$; $w := x + 3$; $v := x$; $u := v + 3$;

After CP: $x := y + z$; $w := x + 3$; $u := x + 3$;

After CSE: $x := y + z$; $w := x + 3$; $u := w$;

Further copy propagation is possible due to the third statement $u := w$.

4.5 Example 2

$a := x ** 2$ $b := 3$

$c := x$

$d := c * c$

$e := b * 2$

$f := a + d$

$g := e + f$

Final form:

$a := x * x$ $f := a + a$ $g := 6 * f$

It is possible for the compiler to get stuck in a "local minima". In the above example, if $f := 2 * a$ was not replaced with $f := a + a$, there could have been possibility to use the shift operator ($f := a << 1$) or to use copy propagation followed by dead code elimination ($a := x * x$; $g := 12 * a$)

4.6 Summary

- Each local optimization does little by itself.
- Often optimizations interact: performing one optimization may enable or disable other optimizations.
- Optimizing compilers can be thought of as a big bag of tricks.

4.7 Typical Structure of an Optimizing Compiler

repeat {apply an optimization rule}

until {no improvement is possible}

- **Convergence** can be guaranteed by defining a metric for performance and ensuring that each iteration improves that metric. This monotonic behaviour would prevent any kind of oscillations that might be possible.
- **Optimality** is not guaranteed. The compiler can get stuck in a local minima.

5 Peephole Optimizations

A peephole optimization is a type of local optimization based on a pattern matching rule consisting of a pattern and a replacement. Here, both pattern and replacement are templates for a sequence of instructions.

$$pattern \rightarrow replacement$$

$$i_1, i_2, \dots, i_n \rightarrow j_1, j_2, \dots, j_m$$

The main idea here is to scan the code and look for the code matching the pattern template and replace it with the replacement code which might be better than the original code in utilization of one of the program's resources. Traditionally the peephole optimization has been quite successfully used on assembly code.

Some examples of peephole optimizations are as follows:

- `move $r1 $r2; move $r2, $r1` \rightarrow `move $r1 $r2`

Possible to do the above replacement because after the first instruction, second move is just a nop and can be removed. It is equivalent to dead code elimination. Do note that the registers mentioned in the instructions are simply placeholders and there can be some other registers as well.

- `addiu $r1, $r1, i; addiu $r1, $r1, i` \rightarrow `addiu $r1, $r1, i + j`

Again registers and constants in the above instructions are simple placeholders and can take any arbitrary value. The above example can be considered as constant folding cast as peephole optimization.

- `addiu $r1, $r2, 0` \rightarrow `move $r1, $r2`
- `move $r1 $r1` \rightarrow `< empty >`

The peephole optimizations are implemented as a database of rules. An optimizer, then scans the code to find the code that matches any pattern in the database and then can be replaced by the replacement.

Many (but not all) local optimizations can be cast as peephole optimizations. Some examples and counterexamples are as follows:

- Algebraic simplification: If an instruction multiplies the value in a register with some power of 2, it can be replaced by a shift instruction.
- Copy Propagation: It cannot be cast as peephole optimization as the definition of a variable and its usage can be very far away with multiple instructions in between with various combinations making it difficult to match them to a pattern

$$w := x$$

...

...

$$t := w + 1$$

Like local optimizations, peephole optimizations can be applied repeatedly. Applying one peephole optimization can activate as well as curb other usages of peephole optimizations. Thus, peephole optimizations are used iteratively a fixed number of times or until no further patterns can be found.

The idea of peephole optimizations becomes more attractive if these can be generated automatically instead of manually coded. Use of enumerative and stochastic methods have been there to automatically learn peephole optimizations.

"Program optimization" is grossly misnamed. Present day "optimizers" have no intention or even pretense of working towards the optimal program. They work towards improving the performance of the code based on certain patterns. Therefore, code "improvers" is a better term.

6 Global Optimizations

(Prepared by Aditya Senthilnathan)

Global optimisations go beyond the scope of a single basic block. It needs reasoning of the whole control flow graph which may be within the body of a method.

6.1 Global Constant Propagation

- This is similar to local constant propagation but now we are reasoning in a global level (across several basic blocks within the body of a method)
- In this optimisation, we basically ask the question "When can we replace the use of a variable **x** with a constant **k**"

Ans: If on **every path** to the use of **x**, if the **last assignment** to **x** is **x:=k**, then we can do this transformation/optimisation.

The algorithm for global constant propagation requires a class of algorithms which are called "Global Dataflow Analysis". This kind of reasoning (with paths, etc.) is required for many different types of optimisation problems in compilers. And so, it helps to create a common framework in which we can implement the same logic but with different kinds of properties so that there is reuse of the same idea.

Let's identify some traits that global optimisations tasks share:

- Optimisation depends on knowing property X at a certain program point P
Eg. In global constant propagation, we were interested in knowing if **x:=k** is the last assignment on all possible paths, at the program point of the last basic block (where **x** is used) in our example graph.
- Proving X at P requires knowledge of entire program
Eg. In our global constant propagation example, to be able to say **x:=k** is the last assignment to **x** on all possible paths that can reach program point **P**, we need to know the characteristics of the entire program. As opposed to local optimisation where you only need info about the concerned basic block. Here, even though transformation is only made in a basic block, it relies on a property that requires knowledge of the entire program.

- It's OK to be conservative

We can either say,

X is definitely true

or say,

We don't know if X is true

We want to get "X is definitely true" as much as possible because only then can we apply the transformation but if we get the conservative answer "We don't know if X is true", then we don't apply the transformation to preserve program correctness. Knowing if X is definitely not true is useless because in this case we still don't apply the transformation and so this use case is clubbed with "We don't know if X is true".

7 DFA Values

(Prepared by Aditya Senthilnathan)

Recall that for global constant propagation, we need the following:

Optimisation: Replace a use of x by a constant k

Constraint: On every path to the use of x, the last assignment to x should be $x := k$

We need to do this for every variable in the program. But first let's do this for a single variable x and then we can generalise it to more variables.

7.1 Dataflow Analysis Values

These are abstract values that are supposed to capture the set or category of concrete values that x (a variable) may take.

At each program point, we associate one of the following values with X (a variable),

Value	Explanation
\top (top)	This value represents either that <ul style="list-style-type: none"> • This statement never executes • or we have not executed or seen it yet
\perp (bottom)	This value represents either that <ul style="list-style-type: none"> • We can't say that X is a constant • X is definitely not a constant
C (constant)	This value represents that X is a constant C

Question: What are we trying to do here? What do these values mean?

Answer: We are trying to figure out what are the possible values of X (a variable) at every program point. The values we are talking about essentially tell us if the program reached here what is the value of X or all the possible values of X at this program point.

Now we will discuss each of the values in a little more detail,

- \top (top) - This value represents that this statement never executes i.e the program/control never reaches here. This might be because this part of the program is never reachable from the beginning of the program. So in this case we just say that the value is \top here because the statement never executes and therefore it is not meaningful to say what this value is.
- C (constant) - It could be any constant. Eg. 0,1,2,-2,etc. This value says that whenever this program point is reached, irrespective of the path taken, we are sure that X will be equal to constant value C at this point.
- \perp (bottom) - This value essentially says that we don't know if X is a constant or not at this point. We are not sure of it's value. This might be happening because it's possible that on different paths X could take different values or on some path X is not a constant, etc. Even if we know for sure that X is not a constant, we still assign \perp to X. This is a conservative assignment. Technically, even if we (DFA algorithm) assign \perp to X at every program point conservatively, then this solution is also correct but it is trivially true and is not of much use to us.

Note: If we're still in the middle of execution of the DFA algorithm and at a particular point, we have value \top for one of the variables in the program then it could also mean that the DFA algorithm hasn't seen this point yet or hasn't reached this point yet. The above discussion and interpretation of \top only applies if a variable has \top value after the DFA algorithm has finished executing

7.2 Using DFA values for transformation

Given global constant information (\top , \perp , C)

- Inspect $x = ?$ (either \top or \perp or C) at the program point that just precedes the statement which you want to examine (the statement that uses x)
- If $x = C$, then replace the use of x with C

For some examples of how to assign DFA values see the lecture video

8 Dataflow Analysis Algorithm (DFA)

(Prepared by Aditya Senthilnathan)

In this section, we will now try to answer the question "How do we identify what the DFA value should be at every program point?"

We will only analyse one particular variable x for now. This can easily be generalised for arbitrary number of variables and this is left to the reader.

8.1 Insight

DFA of a complex and large program can be expressed as a combination of simple rules relating the change in information between adjacent statements.

We're just going to see

- What the value of x is just before this statement
- What is the current statement
- Based on the current statement, we then use rules (to be discussed) to compute what should be the value after the statement

One way to think about this is that we are "pushing" or "transferring" information from one statement to the next. It's almost just like program execution with the caveat that we are not dealing with concrete values anymore but rather abstract DFA values.

8.2 Function for "Constant Information"

We now define a function $C(x, s, in/out)$ where,

- x - variable for which we want to compute abstract values
- s - statement s in the program
- in/out - whether it is input to the statement or output to the statement i.e just before the statement or just after the statement.
 - $C(x, s, in) = \text{Value of } x \text{ just before } s$
 - $C(x, s, out) = \text{Value of } x \text{ just after } s \text{ i.e just after statement } s \text{ is executed}$

The idea is to calculate the value of this function $C(x, s, in/out)$ for every x and every s (both in and out)

8.3 Transfer Function

We need to define how information is transferred from one statement to another i.e how are we pushing information. For this we need to define a transfer function.

Transfer Function: transfers information from one statement to another

We are going to define rules for the following two cases:

- **Case 1:** There are one or more predecessor points with edges to the current program point. In this case, we will see how to combine information/values flowing in from the predecessor points to assign a value for the $C()$ function in the current program point
- **Case 2:** For a given statement s , we see how the information/value is modified as it flows through the statement

8.3.1 Transfer function for Global Constant Propagation

We will define the function by doing an exhaustive case analysis. We can do this kind of finite description of the function with an exhaustive case analysis because there are only a limited number of values and cases to consider.

$C(x, s, p)$ can take only one of 3 different unique values during and after execution of the DFA algorithm i.e \top, \perp, C . Also once the algorithm picks a constant value for x at a particular program point, the constant value is fixed and it cannot vary after that. Eg. If the DFA algorithm picks 3, it cannot assign 4 later on. It can only change to either \top or \perp later on in the execution.

Rules for the Transfer function:

1. If $C(p_i, x, out) = \perp$ for any i , then $C(x, s, in) = \perp$
2. If $C(p_i, x, out) = c$ and $C(p_j, x, out) = d$ and $c \neq d, i \neq j$ for some i, j then, $C(s, x, in) = \perp$
3. If $C(p_i, x, out) = c$ or $C(p_i, x, out) = \top$ for all i , then $C(s, x, in) = c$
4. If $C(p, x_i, out) = \top$ for all i , then $C(s, x, in) = \top$
5. If $C(s, x, in) = \top$, then $C(s, x, out) = \top$
6. If rule 5 does not apply i.e $C(s, x, in) \neq \top$, then $C(x := c, x, out) = c$ [where c is a constant]
7. If rule 5 does not apply i.e $C(s, x, in) \neq \top$, then $C(x := f(...), x, out) = \perp$
8. $C(y := ..., x, out) = C(y := ..., x, in)$ where $y \neq x$

We now have 8 rules relating $C(s, x, in/out)$ to the value of this function at a predecessor point. These rules are exhaustive.

But so far we have just stated the rules. How do we get the actual values?

Think of this as a system of equations. We are interested in solving and identifying values for this C function such that all 10 rules are satisfied.

How do we solve this system of equations?

Note that there is a trivial solution,

$$C(s, x, in/out) = \perp, \text{ for all statements } s$$

All 8 rules are satisfied by this solution but this solution is useless to us because it doesn't allow us to perform any transformations/optimizations. We're not interested in just any solution. We are interested in the "most precise" solution (for some notion of precision). For eg. Saying that $x = 4$ at some program point is more precise than just saying $x = \perp$. Similarly, saying that $x = \top$ is more precise than saying $x = \perp$ because the fact that this program point is unreachable is still useful information because now we can consider doing dead code elimination in this part of the program.

To solve this system of equations in a more precise way, we use a Fixed Point Iteration Algorithm.

Algorithm:


- At every entry point s to the program, set $C(s, x, in) = \perp$
- At every other statement s (which is non-entry), set $C(s, x, in) = \top$
- For every statement s , $C(s, x, out) = \top$
- Repeat the following until all program points satisfy rules 1-8:
 - Pick a statement s not satisfying one or more rules in rules 1-8 and update the corresponding $C()$ function value using the appropriate rule

There is a guarantee that this algorithm will converge. It will also give us the most precise answer. In the worst case, the solution it gives will be \perp at all program points.

9 DFA Example

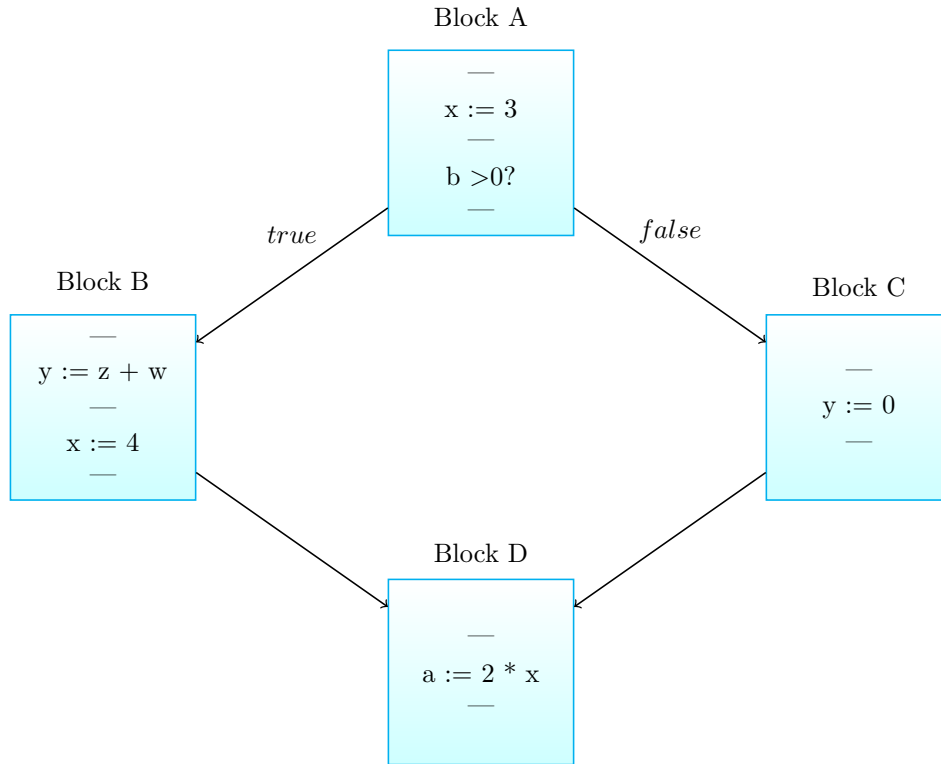
We will apply the fixed point algorithm to the variable “x” in the following piece of code.

Figure 7: Example



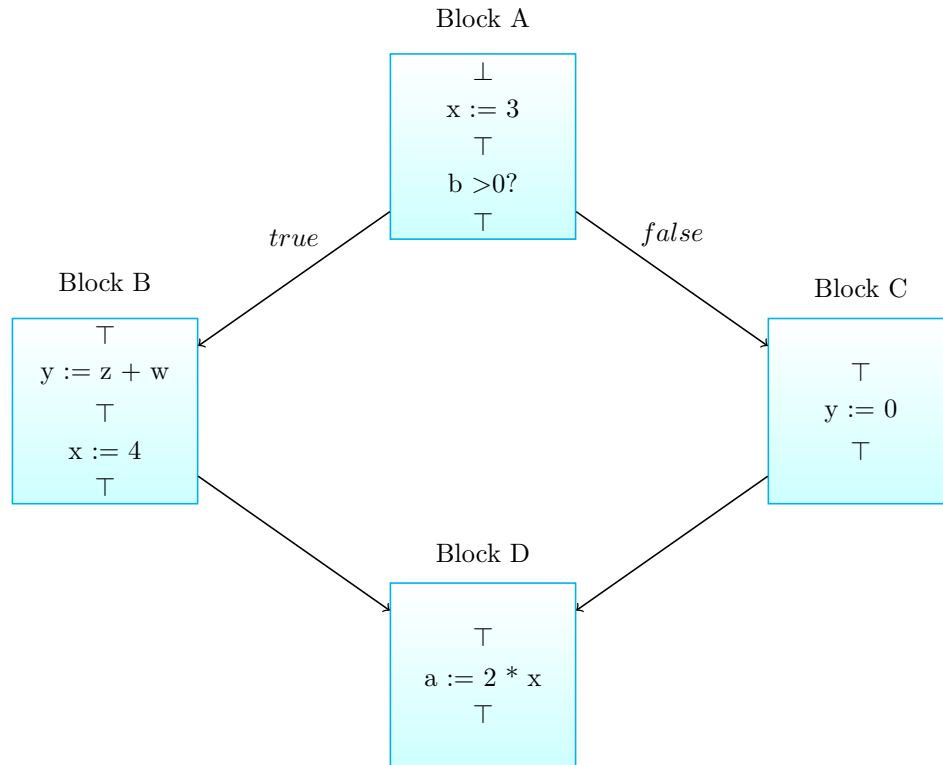
Next figure shows all the program points in our example. Program points are all the points that appear before and after every statement. In case of only one predecessor, the “out of the predecessor” is merged with the “in of the successor” while working on this piece of code. “—” marks the program points in the example we are working on.

Figure 8: Program Points




Algorithm starts by initializing every value to \top (top) except the “in of entry statement” which is initialized to \perp (bottom).

Figure 9: Initialization



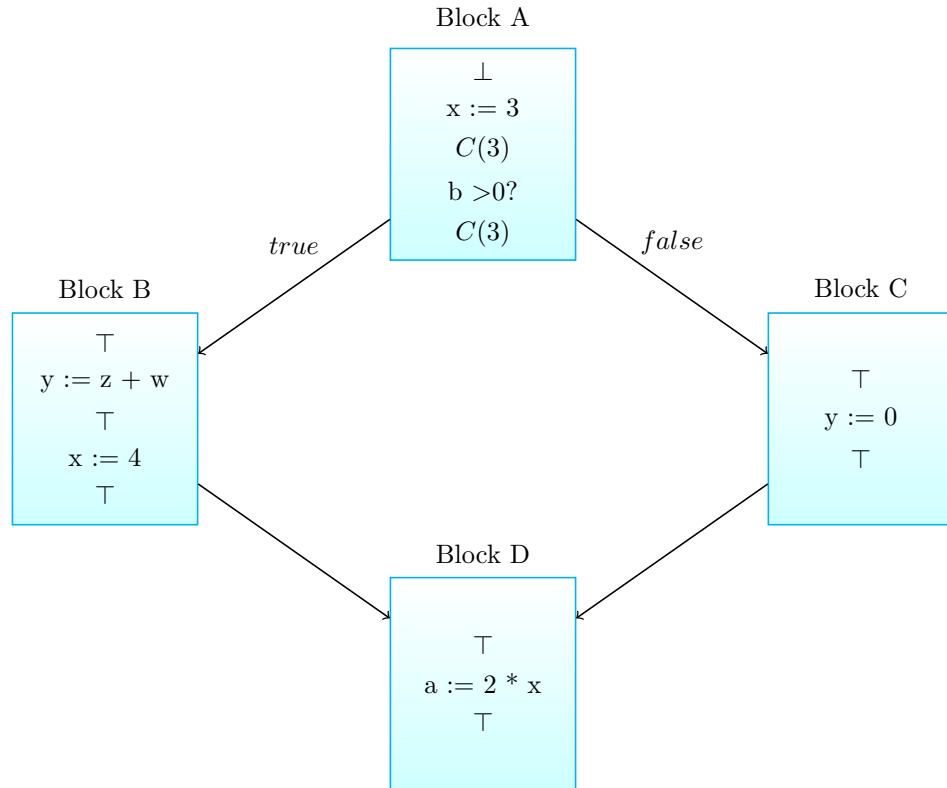
Now the algorithm will pick the statement which does not follow the transfer function rules, which is the “ $x := 3$ ” statement in this example, and update it. The “out of the statement” should be Constant 3 according to the transfer function rule. After updating the “out of the statement”, we get this graph.

Figure 10: “ $x := 3$ ” statement update



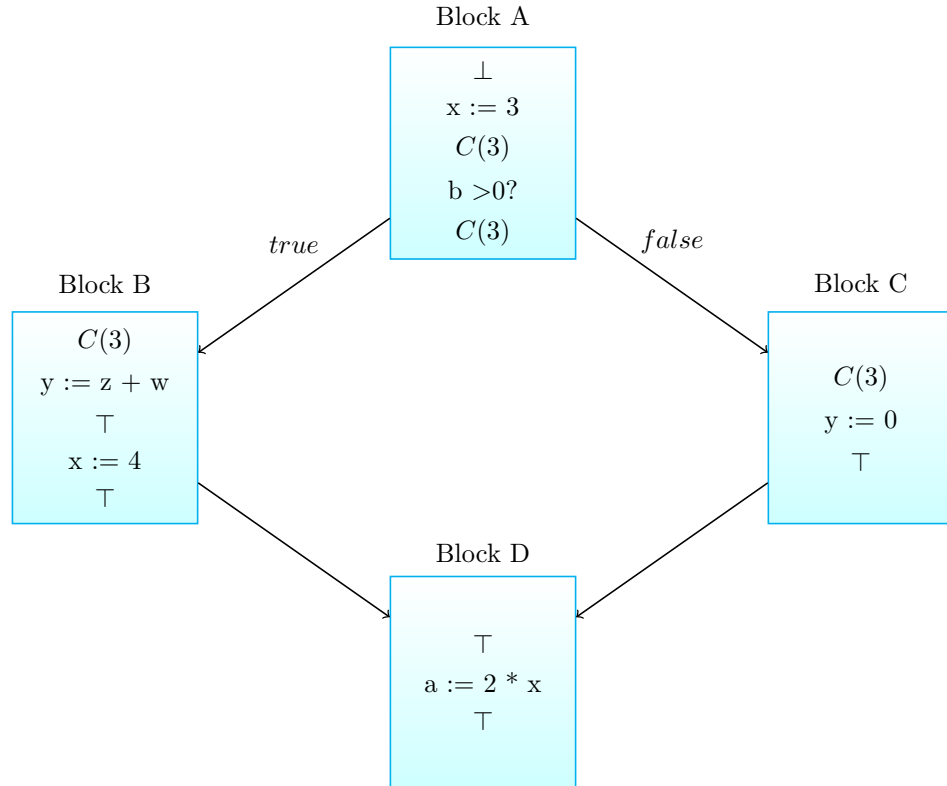
The algorithm repeats this process until all program points satisfy the transfer function rules.

Figure 11: “b > 0?” statement update




At this stage, there are two program points which do not satisfy the transfer function rules. The entry point of Block B and Block C do not match with their predecessors. We can update them.

Figure 12: “Successors of Block B and Block C” update



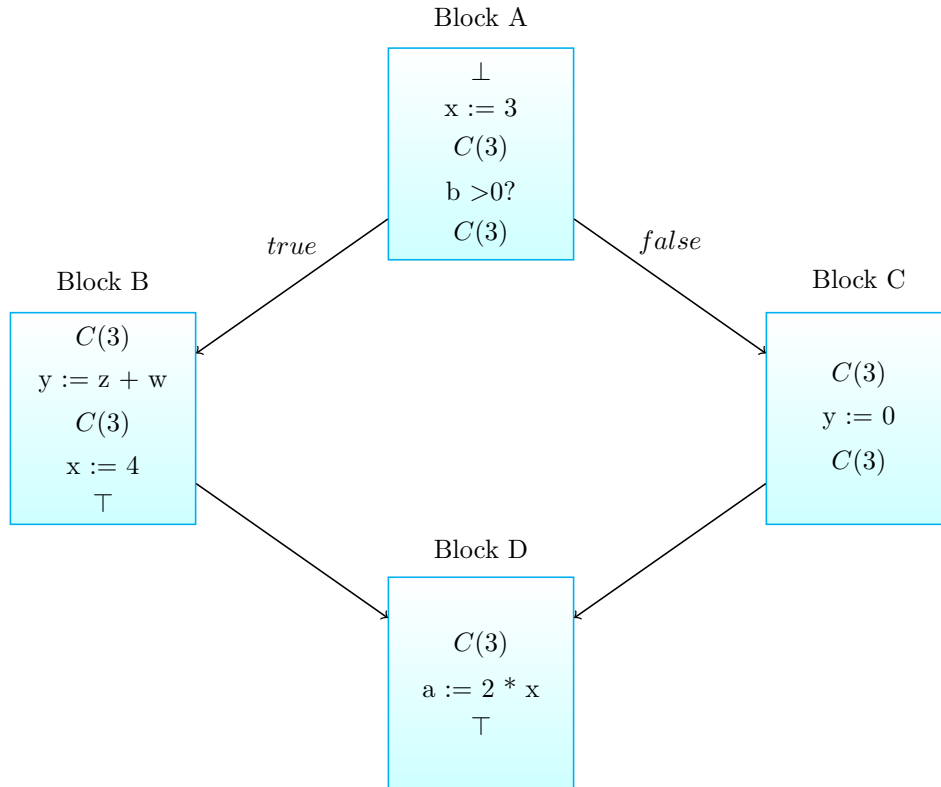
Again there are two different statements that do not satisfy the transfer function rules, we can update them.

Figure 13: “ $y := z + w$ ” and “ $y := 0$ ” statements update




At this stage, we can either update “ $x := 4$ ” statement, or we can update the successor of Block D. Irrespective of which one we pick, algorithm will return the same answer. In this example, we will update the successor of Block D. One predecessor is a constant, and other predecessor of Block D is \top . Using the 3rd transfer function rule, we can update the value of successor to the same constant.

Figure 14: “Successor of Block D” update




Updating “a := 2 * x” statement in Block D.

Figure 15: “a := 2 * x” update




There is only one statement, “x:=4” in Block B, that does not satisfy the transfer function rules. After updating it, we get this graph.

Figure 16: “x := 4” statement update



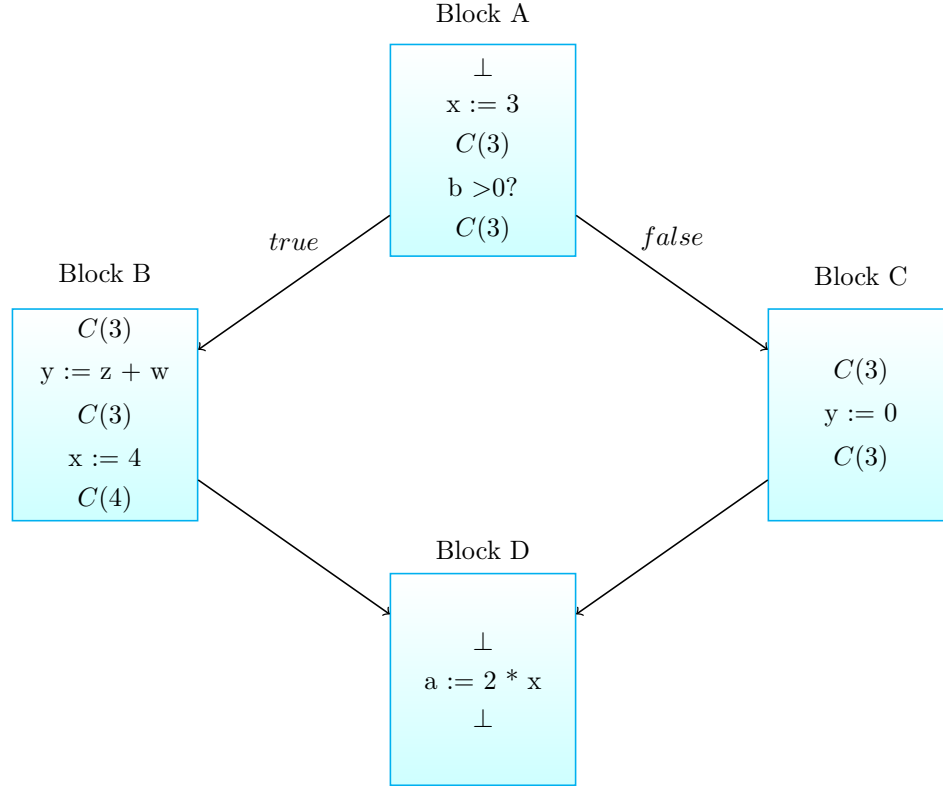
Because of the previous update, the successor of the Block D does not satisfy the transfer function rules. According to the 2nd transfer function rule, it should have the value \perp .

Figure 17: “Successor of Block D” update



Updating the “ $a = 2 * x$ ” statement in Block D, we get this graph.

Figure 18: “ $a := 2 * x$ ” statement update



At this stage, there is no program point that violates any of the 8 transfer function rules, so the algorithm terminates by returning this **Fixed point solution** to us.


9.1 Some Guarantees

- The fixed point solution will satisfy all the transfer function rules at each program point.
- The algorithm is guaranteed to converge, irrespective of whether the program has loops or not.
- Whenever there are multiple options to choose for a statement S to be updated, irrespective of which one we choose, we always get the same answer.
- The fixed point solution is going to have some precision guarantee, by some definition of precision, for the system of equations we are trying to solve.

10 DFA Loop Example


In this section, we will look at an example of DFA for a program with loops. We will apply the DFA algorithm on the variable “x” in the following piece of code.

Figure 19: Example




Next figure shows all the program points in our example. Program points are all the points that appear before and after every statement. In case of only one predecessor, the “out of the predecessor” is merged with the “in of the successor” while working on this piece of code. “—” marks the program points in the example we are working on.

Figure 20: Program Points




Next step is to initialize all the program points. To do that, we initialize all the program points to value \top except the entry points of the program which are initialized to \perp

Figure 21: Initialization




In the next step, we identify the statements which do not follow the transfer function rules and start updating them. At this stage, only the statement “ $x := 3$ ” does not follow the transfer function rule, so we update it.

Figure 22: “ $x := 3$ ” statement update




The process of identifying the statement that violates the transfer function rules and updating it, is repeated until every program point satisfies them. At this stage, the “b > 0?” statement violates the transfer function rules, so we update it.

Figure 23: “b > 0?” statement update




At this stage, the successors of Block B and Block C do not satisfy the transfer function rules, so we can update them.

Figure 24: “Successors of Block B and Block C” update




There are again two different statements “ $y := z + w$ ” and “ $y := 0$ ” that violate the transfer function rules, and we will update them.

Figure 25: “ $y := z + w$ ” and “ $y := 0$ ” statements update




We have again two different statements that we can update at this stage, and we can choose to update anyone that we like, since the algorithm is guaranteed to return the same solution. In this example, we will update the “ $x := 4$ ” statement.

Figure 26: “ $x := 4$ ” statement update




Now we will do the same thing for Block D, update the successor for Block D and then update its statements one by one. At the end, we will reach this stage.

Figure 27: “Block D” update




At this stage, we haven't reached the fixed point solution. The successor of Block C violates the transfer function rules. Earlier, it was acceptable for the algorithm to update the value to $C(3)$ because the predecessor entering from Block D was \top , which is now updated to \perp . So, we will have to update this program point.

Figure 28: "Successor of Block C" update



Updating the “ $y := 0$ ” statement, we reach this stage.


Figure 29: “ $y := 0$ ” statement update



All the program points at this stage satisfy the transfer function rules, so this is a **Fixed Point Solution**.


To show how the algorithm could have worked differently, we will change the “ $x := 4$ ” statement to “ $x := 3$ ” in Block C of our same example. This is our new example now.

Figure 30: Example




Algorithm will again start by initializing everything to \top , except the entry statement which is initialized to \perp . We will follow the same procedure and eventually reach this point after updating Blocks A, B and C.

Figure 31: “Blocks A, B and C” update



This time the successor of Block D will be updated to constant 3 instead of \perp , because both the predecessors have the same value of constant 3. Updating the successor, and the following instructions in Block D, we reach this stage.

Figure 32: “Block D” update



Question : Have we reached the fixed point now?

Answer : It turns out that, in this case, we have reached the **Fixed Point Solution**. Both the predecessors of Block C, coming from Block A and Block D, have the value of Constant 3, so there is no need to update it.

11 DFA Value Orderings

In this section, we are going to discuss the convergence property of the DFA algorithm. And we will start by defining some arguments for it.

11.1 Orderings

Idea: Simplify the presentation of analysis by ordering the abstract values. We will create an arbitrary operator, “<” (less than), which orders them in the following manner.

$$\perp(\text{bottom}) < C(\text{constant}) < \top(\text{top})$$

This operator is a transitive which means that $\perp(\text{bottom}) < \top(\text{top})$ is also true.

We are going to use this operator to reason about the convergence properties of the DFA algorithm.


11.2 Partial Ordering

This ordering that we defined is a partial ordering, and not a total ordering. It means that among all the possible values, it is not necessary that all values are comparable. There are some values that are comparable, and some values that are not.

Another way to represent the partial ordering is by using a vertical representation instead of a horizontal representation, and using the directed arrow to represent the “<” (less than) operator. The arrow from \top to C indicates that $C < \top$.



C is not a value, but a place holder for any constant. A better way to represent the ordering would be in the following manner.



Notice that the figure does not have any relation between 0 and 1, or 1 and 2, or between any two constants, in general. That is why it is called a partial ordering, because only some values are related by “<” operator.

Greatest Value: The value which is not less than any other value. In this case, it is \top (top).

Least Value: The value which is not greater than any other value. In this case, it is \perp (bottom).

11.3 Greatest Lower Bound

To build the concept of greatest lower bound, we introduce " \leq " (less than equal to) operator which in addition to all the properties of " $<$ " (less than) operator, is also reflexive, *i.e.*, $x \leq x \forall x$.

The **greatest lower bound(glb)** of x_1, x_2, \dots, x_n is the greatest value that is lower than (by " \leq " operator) $x_i, \forall i$ s.t. $1 \leq i \leq n$.

Here are some examples of glb:

1. $\text{glb}(\top, 1) = 1$
2. $\text{glb}(\top, \perp) = \perp$
3. $\text{glb}(2, \perp) = \perp$
4. $\text{glb}(1, 2, \top) = \perp$
5. $\text{glb}(1, 2) = \perp$

Observation: Glb can be used to replace the first four transfer function rules which defined the relation between the "out of the predecessors" and "in of the successor".

$$C(s, x, in) = \text{glb}\{C(p, x, out) \mid p \text{ is a predecessor of } s\}$$

11.4 Convergence Argument

The DFA algorithm repeats itself until nothing changes. The only reason for it to not converge is, if it keeps changing forever. There are two reasons for it to keep changing forever.

1. Algorithm can keep taking a different value from the infinite set of values, for the same variable and at the same program point.
2. Algorithm can keep oscillating between the finite number of same values, for the same variable at the same program point.

Using glb, we will prove that the fix point iteration always converges. Here are the convergence arguments:

1. Values start at \top and can only decrease.
2. \top can change to C , and C to \perp .

It cannot happen that \top changes to some constant C , let's say 1 and in a different iteration it changes to a different constant, let's say 0. This is because one of the predecessors would still have the value 1, and if any other predecessor changes its value to a different constant, we are going to take the glb of predecessors which only allows the value to decrease in the order mentioned in this argument. This intuitive argument can be proved more formally by using induction on the Control Flow Graph of the program.

Since this argument does not allow the value to go upwards, *i.e.*, from C to \top , or from \perp to C , or from \perp to \top , oscillation is not possible among these values.

3. Thus, $C(s, x, in/out)$ can change at most twice, $\forall s$ and x .

11.5 Worst Case Execution Time

11.5.1 For One variable

At least one value changes in each iteration, and at every program point a value can change twice therefore:

Number of steps of Fixed Point Iteration Algorithm \leq (Number of C(s, x, in/out) statements) * 2

There can be at most two C(s, x, in/out) statements per each statement therefore:

Number of steps of Fixed Point Iteration Algorithm \leq (Number of program statements) * 4

11.5.2 For all variables

There are two options to deal with all the variables.

1. Run the algorithm separately for each variable.
To calculate the worst execution time, we multiply the number of variables to worst execution time for one variable. The number of variables would be less than the number of program statements, assuming Three Address Code. Therefore, the worst case would be quadratic in size of the program.
2. Keep track of all variables simultaneously.
This can be done by modifying the transfer function such that it looks at every variable simultaneously and updates multiple values in one step. It improves the performance because looking at a set of variables is usually cheaper than looking at each of them separately. It also provides more information while updating the values. For example, if we have a statement like “ $x := y + z$ ”, and we have the information that y and z are constants at this point, then we can also infer that x is a constant. Had we dealt with the variables separately, this would have not been possible. In another words, the analysis would have been weaker. If we do things separately, we may end up with less precise solution in some situations, for some type of analysis.

12 Feb 19 Discussion

- **Jai Javeria:** When discussing global constant propagation, it was said that to apply this transformation, we need to know that property X applies at a particular point and to know that this property X applies at this point we need knowledge of the entire program. However, this is not always true and is not a necessary condition and so it is more accurate to say that to know that this property X applies at this point we *may* need knowledge of the entire program.

- **Sonu Mehta:** How do we apply optimizations? Do we apply local optimizations first and then global?

A: Global optimizations subsume local optimizations. So, typically we'll be more concerned with global optimizations than local optimizations because of the more general nature of global optimizations. Most of the initial optimizations will all be global. However, there are specialized local scenarios where local optimizations come in handy and are helpful. So at the end of the optimization pipeline we apply the local optimizations. These local optimizations are typically difficult to do globally.

- **Shubham Sondhi:** In module 77, it was said that assigning bottom to all program points would be a valid solution but it is actually not a valid solution because it would violate the rule which says that if we do not see a value \top in the "in" of a statement, and the statement is $x:=c$, then we should assign a constant to the out of the statement. Everything \top is also not an answer because it doesn't satisfy the boundary condition that the in of the beginning condition must be \perp . Here, what might work is slightly modifying the rule being violated to say that if "in" of a statement is not \top and the statement is $x:=c$, then the out can be c or lower (which is \perp). With this rule, all \perp will be a valid solution.

- **Anirudh Panigrahi:** In rule 7, it was said that if a statement is $x:=f(\dots)$ and in of this statement is not top, then the out of this statement will be bottom. Is the RHS of this statement necessarily a function or can it be anything which is not a constant?

A: It can be anything which is not a constant. The $f(\dots)$ is just a placeholder for anything not a constant.

- **Jai Javeria:** It was said in Module 78 that the DFA values have a partial order but it would be more accurate to say that its a strict partial order because partial orders can be reflexive.

- **Jai Javeria:** Dataflow was discussed in the context of tracking one variable's values across different program points. It was mentioned that handling multiple variables might be potentially faster than handling just single variables?

A: There is a time vs. space tradeoff. If we have a table/set of variables storing each variables values at different program points, we could maintain that table with efficient data structures. Details will be clearer when you watch later modules.

- **Sonu Mehta:** Dataflow analysis was discussed in relation to global constant propagation. Can similar things be done for other optimizations as well.

A: Yes. Dataflow analysis is a general framework that can be applied to many different frameworks.

- **Sonu Mehta:** The greatest lower bound (glb) was discussed and it was said that it was transitive and reflexive. Is $\text{glb}(1,1)=1$?

A: Yes. It is reflexive.

- **Indrajit Banerjee:** Instead of just having 3 values of \top , \perp and constant, is it possible to extend this analysis by including a new abstract value which tells whether x is definitely greater than 0 or we cannot say anything about it, at some program point?

A: We need to get some new transfer function rules, for the different kind of statements that would affect the variable in consideration.

- **Arpit Saxena:** Wouldn't it be difficult to define the transfer function rules for a variable to be definitely positive at a program point? Even for a simple statement like $x = x + 1$, we cannot surely say that the "out of the statement" would be greater than 0, given that the "in of the statement" was positive, because the value might overflow in some case.

A: It is tied to the language semantics. A language like C considers interger overflow an undefined behaviour. We can ignore the overflow cases while defining the transfer function rules for a language like C. For other languages, We can maintain a set of potential values that x could have a program point, every time we find a new value that x can take, we add it to the set. But this would take a lot of memory and time, in some cases it may not even converge.

- **Jai Javeria:** Once we convert the language to its intermediate representation like LLVM, do we lose information for example, whether an integer can overflow or not?

A: LLVM preserves information in some cases(for example, whether integer overflow is an undefined behaviour or not), and in other cases it choses the most conservative answer.

- **Vijay Bhardwaj:** How is the dataflow technique applied in the multi-threaded programs, or in case of pointers?

A: For multi threaded programs, we make an assumption that there are no data races, even Java and C makes this assumption. Data races in C are undefined behaviour, that's what allows the compilers to do all the optimizations.

Nothing changes in case of Pointers. Pointer analysis and Aliasing analysis(whether two pointers are pointing to the same location or not) is also modelled as dataflow analysis.

- **Aditya Senthilnathan:** If C assumes that data races are undefined behaviour, does the compiler fail silently or detect such data races and warn user about it?

A: Compiler typically fails silently. Detecting the data races is very expensive and it would make the program very slow, so it is not done in production code.